

# Core Labs



### Deep Learning Best Practices on Ibex

GLENDON HOLST, HPC ML & VISUALIZATION DAVID PUGH, DATA SCIENTIST



### Deep Learning on Ibex

- Quick Intro
  - Prerequisites: Ibex, Slurm, Conda, ...
- Deep Learning Best Practices
  - Get More Done with Available Resources
  - Get Faster Access to More Resources
- Hands-On Clinic By Appointment

### **Ibex GPU Composition**



GPU Model	#	Mem	СРИ	Cores	Cores per GPU	Memory	Mem per GPU	Storage	Nodes	Total
GTX1080Ti	4	11.5GB	Haswell	36	9	256 GB	64 GB	200GB	8	32
GTX1080Ti	8	11.5GB	Skylake	32	4	384 GB	48 GB	900GB	4	32
Tesla P100	4	16 GB	Haswell	36	9	256 GB	64 GB	200GB	6	24
Quadro P6000	2	22 GB	Haswell	36	9	256 GB	64 GB	200GB	2	4
RTX2080Ti	8	11 GB	Skylake	32	4	384 GB	48 GB	900GB	4	32
Tesla V100	4	32 GB	Skylake	32	8	384 GB	48 GB	800GB	8	32
Tesla V100	8	32 GB	Cascade Lake	48	6	768 GB	96 GB	32TB	30	240

### Getting Started: Workstation → Ibex

### Workstation

Personal – all the resources!

- Interactive
- Local Storage
- Packaged Software
  - Conda Envs

#### • Ibex

/ibex\_

/scratch

- Shared specify resources; CPU, Mem...
- Batch scripted
- Remote Filesystem
- App Stack Modules

Conda Envs

### Getting Started: Workstation → Ibex

- ssh vlogin.ibex.kaust.edu.sa
- First login auto-generates keys & ssh config
  - .ssh/config



https://www.hpc.kaust.edu.sa/ibex/new\_user https://www.hpc.kaust.edu.sa/ibex/faq

### Batch Training: Workstation → Ibex

### WS: Parameterize train.py

### WS: Write shell script to perform training

```
- #!/bin/bash
conda activate ./env
python train.py "$@"
```

```
- ./run_train.sh --batch-size=64
```

### Batch Training: Workstation → Ibex

### Ibex: Add #SBATCH resource specs

```
- #!/bin/bash
  #SBATCH --nodes=1
  #SBATCH --ntasks=1
  #SBATCH --mem=45G
  #SBATCH --cpus-per-task=4
  #SBATCH --gres=gpu:1
  #SBATCH --constraint="[p100|gtx1080ti]&intel"
   #SBATCH --time=24:00:00
  #SBATCH --partition=batch
   #SBATCH --output=log-%x-slurm-%j.out
   #SBATCH --error=log-%x-slurm-%j.err
   . . .
```

### Batch Training: Workstation → Ibex

## Ibex: Run training job via Slurm

- sbatch run\_train.sh --batch-size=64
- sbatch --time=15:00 --constraint=[p6000] --partition=debug \
   run\_train.sh --batch-size=16

## Ibex: Slurm job management

- squeue -u \$USER
- scancel <jobid>
- ginfo --all
- sinfo -o "%20N %10c %10m %25G %f" | grep -E '(FEATURES|gpu:)'
- scontrol show job <jobid>

### Deep Learning: Best Practices on Ibex

- Get More Done with Available Resources
  - Making your jobs run faster helps everyone

- Get Faster Access to More Resources
  - Good things come to those who share...

#### Best Practices: Allocate Sufficient...

### Allocate Sufficient CPUs

- CPU\_PER\_GPU
  - Depends on node configuration (Ibex hardware chart)
  - Specify 4, or proportional fair share (from chart)
- --cpus-per-task=\$(( NUM\_GPUS \* CPU\_PER\_GPU ))

## Allocate Sufficient Memory

- MEM\_PER\_GPU in GB
  - Depends on node configuration (Ibex hardware chart)
  - Specify 45, or a little less than proportional fair share (from chart)
- --mem=\$(( NUM\_GPUS \* MEM\_PER\_GPU ))G

## Load Training Data from Local Storage

- /tmp
- AI GPU nodes have very large and fast local storage:
  - Less IO contention with other users; fast SSD
  - Only visible on the compute node it is attached to (per node + job)
  - Local storage is temporary only lasts as long as the job

#### Use Reference Datasets

- Eliminates initial dataset copy time
- Faster local storage
- Use sinfo command to find nodes with reference data:
  - sinfo -o "%20N %10c %10m %25G %f" | \
     grep -E '(AVAIL\_FEATURES|ref\_)'
- Ensure GPU, CPU, Memory, and GRES match job spec; then add constraint:
  - --constraint=[ref\_32T]
- Provided reference data is available on node local storage:
  - /local/reference/<reference-data-set>
- Request to provide dataset: <ibex@hpc.kaust.edu.sa>
  - /ibex/reference/CV/{COCO, GQA, ILSVR}

## Copy Training Data to Local Storage

- Ensure requested nodes have sufficient local storage (Ibex hardware chart)
- For single node allocation:
  - cp -r /ibex/scratch/\${USER}/data/set /tmp/data
  - cp /ibex/scratch/\${USER}/data/set.tgz /tmp/data cd /tmp/data; tar xvpf set.tgz
- For single or multi-node allocations:
  - sbcast /ibex/scratch/\${USER}/data/set.tgz /tmp/data
     cd /tmp/data ; srun -N \$N -n \$N tar xvpf set.tgz
  - Broadcast copy operation to all nodes

## Filesystem Guidance

- For improved copy performance, prefer source data with fewer, larger files
  - Avoid using large directories with many (e.g., thousands of) small files, especially on /ibex/scratch/
  - Archive source data in \*.tar or \*.zip files (and un-compress on the compute node after the copy)
  - \* Or, place individual files inside containers, like an HDF5 file, and work with these files via the container API
- Avoid using /home for data files (read or write)

### Best Practices: Loading Data...

## Maximize data loading performance

- Use framework provided data loading and processing tools:
  - E.g., tf.io.\*, tf.image.\*, tf.data.Dataset.\*, tf.data.\*.AUTOTUNE
- Load and process data in parallel
  - E.g., .map(preprocess, num\_parallel\_calls=AUTOTUNE)
- Ensure that data buffers are sufficiently large (for hiding IO latency and caching)
  - .shuffle(args.shuffle\_buffer\_size, reshuffle\_each\_iteration=True)
     .prefetch(args.prefetch\_buffer\_size)
- See example:
  - https://github.com/kaust-vislab/tensorflow-gpu-data-science-project
    - src/horovod-keras-example/train.py

## Maximize utilization of large-mem GPUs

- E.g., Ibex V100 has 32GB (2x standard 16GB V100) use it
- Double batch size, double learning rate, half number of GPUs per job
  - Job performance stays the same (similar)
  - Double number of simultaneous jobs
- Deep Learning at Scale, Large Mini-batch SGD
  - https://towardsdatascience.com/deep-learning-at-scale-accurate-large-mini-batch-sgd-8207d54bfe02
  - Linear Scaling Rule: Ir' = Ir \* k (for k time batch size increase)
  - Warmup Strategies:  $lr \rightarrow lr'$ , gradual constant increase from small lr over 5 epochs
  - Batch Normalization: Variety of techniques... subtle pitfalls\*
- \* Ask <ibex@hpc.kaust.edu.sa> for assistance with scaling on Ibex

### Monitor GPU utilization

- Even basic GPU and CPU monitoring can diagnose common performance issues

## Add monitoring to job script

## Interactively monitor running job

```
- # show running jobs and find <jobid>
squeue -u $USER
```

```
- # examples of interactive session GPU / system monitoring
> srun --jobid=<jobid> -u --pty bash -i
> srun --jobid=<jobid> -u --pty nvidia-smi dmon
> srun --jobid=<jobid> -u --pty top -H -u $USER
```

## GPU compute utilization can exhibit:

- Consistent (e.g., > 90%) compute utilization this is good
- Consistent middling (e.g., < 70%) utilization let's improve</li>
  - Small batch size? Low profile sampling rates (hidden oscillation)?
- Multi-GPU request; some GPUs have 0% utilization let's improve
  - GPU device not bound process in nvidia-smi pmon?
    - Initialization issue
  - GPU devices are bound to process, but show 0% utilization?
    - Framework need explicit multi-gpu support added to code
    - See Utilize all allocated GPUs

## GPU compute utilization can exhibit:

- Oscillation between high (e.g., 100%) and low (e.g., 10%) utilization let's improve
  - Insufficient CPUs to load / process batch data?
    - Ensure sufficient CPUs per GPU, and framework data loader uses them
  - Data load is not in parallel with training operation?
    - Use framework data loader utilities with sufficient CPUs
  - GPU isn't busy for long enough?
    - Increase batch size and buffer sizes to help hide latency
  - Slow filesystem IO performance?
    - Load data from local storage
  - Checkpointing pauses?
    - Save checkpoint to local storage; copy to /ibex/scratch concurrently

#### Utilize all allocated GPUs

- Multi-GPU support requires support in / changes to training code...
- TensorFlow: Multi-GPU aware; use tf.distribute.\*Strategy
  - strategy = tf.distribute.MirroredStrategy()

```
BATCH_SIZE = 64
BATCH_SIZE = BATCH_SIZE * strategy.num_replicas_in_sync
with strategy.scope():
   model = create_model()
```

- PyTorch: Tensors are assigned to GPU devices manually
- Horovod: Distributed TensorFlow / PyTorch training over MPI
  - Ask <ibex@hpc.kaust.edu.sa> for assistance with scaling on Ibex

### Best Practices: Writing Models...

## Checkpoint to local storage

- Local storage (/tmp) is fast...
- Local storage is temporary!
- Write (rsync) to persistent storage (/ibex/scratch) in background

### Best Practices: Writing Models...

## Copy in background...

```
# configuration
   RSYNC DELAY SECONDS=3600
   PID CHECK DELAY SECONDS=60
   RSYNC_DELAY_LOOP=$((RSYNC_DELAY_SECONDS / PID_CHECK_DELAY_SECONDS))

    # launch training command in background, get process ID on next line

   python train.py "${LOCAL_CKPNT_DIR}" ${TRAINING_ARGS} &
   RUN PID=$!
   # asynchronous rsync of training logs to persistent storage
   RUN DONE=false
   while [ "${RUN_DONE}" != true ] ; do
     for i in $(seq 1 ${RSYNC_DELAY_LOOP}); do
       sleep ${PID_CHECK_DELAY_SECONDS}
       if [[ "$(ps -h --pid $RUN_PID -o state | head -n 1)" = "" ]] ; then
         RUN DONE=true ; break
       fi
     done
     rsync -a "${LOCAL_CKPNT_DIR}/" "${PERSISTENT_CKPNT_DIR}"
   done
```

### Deep Learning: Best Practices on Ibex

- Get More Done with Available Resources
  - Making your jobs run faster helps everyone

Get Faster Access to More Resources



Good things come to those who don't wait...

### Bonus Resources: gpu\_wide24

## New gpu\_wide24 partition



- For wide (>= 4 GPUs) and short (< 24 hrs) jobs
- 4 AI-GPU nodes with 8xV100 GPUs connected via NVLink
- Benefits:
  - Faster turn-around for early results
  - Interleaved jobs make regular progress
  - Access extra resources
- Satisfying jobs are automatically eligible to run in the partition
  - E.g., --time=24:00:00 --gres=gpu:v100:4

### Best Practices: Checkpoint Support...

## Add checkpoint / restore support

- Checkpoints are written / read by single task usually root rank restored to all
  - For single-task multi-GPU, <checkpoint> and <restore> are sufficient
  - E.g., For Horovod
- Checkpoint to local storage (/tmp); copy to persistent storage in parallel
- Restore from persistent storage
- Checkpoint at reasonable intervals regular, but not too frequently
  - checkpoint delay should be small fraction of epoch training time

### Best Practices: Checkpoint Support...

## Write checkpoint to local storage

**Note**: If save\_freq = int, it is the number of samples (not batches) to process before saving

### Best Practices: Checkpoint Support...

### Restore checkpoint from /ibex/scratch

```
E.g., TensorFlow
- ckpnt dir = pathlib.Path(args.read ckpnt from)
   ckpnt_filepath = None
   initial epoch = 0
   for _epoch in range(args.epochs, 0, -1):
       _ckpnt_filepath = f"{ckpnt_dir}/ckpnt-epoch-{_epoch:02d}.h5"
       if os.path.exists(_ckpnt_filepath):
           ckpnt_filepath = _ckpnt_filepath
           initial_epoch = _epoch
           break
   if ckpnt filepath is not None:
     model_fn.load_weights(ckpnt_filepath)
```

### Best Practices: Sequence Jobs...

## Split training over multiple jobs

- Shorter job runtime means more jobs to run; each job does a fraction of the work
- Modify training code
  - Restore from latest checkpoint file
  - Process a fixed number of epochs / fixed time limit...
  - Exit cleanly (before job terminates).
  - Use reference data on local storage
- Launch multiple jobs at once
  - Jobs must run in sequence
    - i.e., Next job depends upon previous job successfully completing
  - Automate (script) launches via Slurm sbatch
  - Future: Use workflow management tools (e.g., decimate) support in development...
    - Improved: error retry, early exit, management tools

#### Best Practices: Launch Jobs...

### Create launch\_jobs.sh script

- #!/bin/bash

#### Best Practices: Launch Jobs...

## Create launch\_jobs.sh script

Launch all jobs at once; each job depends upon the previous job completing successfully

## Manage jobs manually

- squeue -u \$USER
- scancel <*jobid*>

#### Best Practices: Clean Exit...

## Exit cleanly prior to job termination

- Do not wait until job gets killed
  - Termination can cause checkpoint corruption (if it interrupts the copy operation)
- Ensure ample time provided to job to complete work given (e.g., number of epochs)
- Notify yourself in cases where jobs fail, timeout, or come close to running out of time:

```
#SBATCH --mailtype=FAIL,TIMEOUT,TIMEOUT_90
#SBATCH --mailuser=<username@kaust.edu.sa>
```

### Use reference datasets

Eliminate dataset copy times — important for shorter jobs

### Best Practices: Wrapping Up...

## Engage Us

- gpu\_wide24 & ref\_32T early adopters provide feedback, ask questions
  - Queue wait times? Job overhead? Experience? Scaling?
- Contact us:
  - Email: <ibex@hpc.kaust.edu.sa>
  - Slack: https://kaust-ibex.slack.com #general, #gpu, #conda

## Example Projects

- https://github.com/kaust-vislab/tensorflow-gpu-data-science-project
  - src/horovod-keras-example/train.py
- https://github.com/kaust-vislab/pytorch-gpu-data-science-project
  - src/horovod-example/train.py